

Category	Name	Description	* If you wish to contribute anything, please leave comments in this sheet - we'll verify the resource and add it accordingly. Help out your fellow classmates!						
Go through first	<a href="#">State of GPT</a>	A talk by Andrej Karpathy to get you up to speed on LLMs							
	<a href="#">A Hackers' Guide to Language Models</a>	A 90 minute crash course on language models, using APIs and open-source models, and finetuning							
	<a href="#">okay but I want GPT to perform 10x for my specific use case</a>	A short video on how to plan for improving your model's performance for a specific use case							
	<a href="#">Multimodal 101 Tutorial</a>	Exploring the use of Multimodal Models to tackle tasks that require jointly using text and image							
	<a href="#">Intro to Large Language Models</a>	Another talk by Andrej Karpathy on topics surrounding LLMs							
Quick Courses	<a href="#">LLM Course - Maxime Labonne</a>	The most popular resource to learn LLMs Offers roadmaps and content on every aspect of using these models Highly recommended to skim through to find what's next to learn							
	<a href="#">deeplearning.ai Short Courses</a>	Quick 1-2 hour courses on surface-level topics relating to recent developments Some topics include: Prompt Engineering, LangChain, Finetuning, Diffusion Models							
	<a href="#">W&amp;B Courses</a>	More comprehensive but slightly longer tutorials							
	<a href="#">HuggingFace NLP Course</a>	Very good for getting an overview of general topics in NLP The best for getting a hang of using transformers Look through the blogs for newer developments							
	<a href="#">LangChain Playlist - DataIndy</a>	The most popular YouTube resource for learning LangChain Slightly outdated							
	<a href="#">Made With ML</a>	Covers aspects of ML from a development/deployment POV Less relevant to this course but still worthwhile to explore							
Catalogs	<a href="#">awesome-langchain</a>	Collection of repositories of LangChain projects Highly recommended to go through this for inspiration!							
	<a href="#">State of LLM Apps (Streamlit)</a>	Information, statistics, a gallery and more. All related to applications made using LLMs. Provided by streamlit.							
	<a href="#">rentry Local Models Related Links</a>	Huge catalog of links related to papers, tools, datasets, leaderboards etc. Covers domains outside of just text: Vision, Speech, Visual-Language etc.							
	<a href="#">open-TTS</a>	A repo tracking open-source TTS models as they come out (could help with ASR too)							
	<a href="#">LangChain Tutorials</a>	Quick tutorials from the official LangChain site							
	<a href="#">Weaviate Tutorials</a>	Quick tutorials from the official Weaviate site							
Tools/Docs	<a href="#">LangChain</a>	Popular framework for working with language models in general							
	<a href="#">OpenAI API</a>	Endpoints for inference using proprietary models like GPT-4, DALL-E 3, and their Embedding engines							
	<a href="#">Anyscale</a>	Popular platform for developing, deploying and scaling AI applications							
	<a href="#">HuggingFace Models</a>	Collection of open-source models across a variety of domains <b>LOOK FOR MODELS OUTSIDE OF JUST LLMS HERE (Multimodal, Vision, Speech etc.)</b>							
	<a href="#">Diffusers</a>	The go-to library for running pretrained Diffusion models (for text-to-image generation)							
	<a href="#">axolotl</a>	Popular minimalistic Framework for LLM finetuning							
	<a href="#">ollama</a>	Framework for LLM inference							
	<a href="#">unsloth</a>	ollama but cooler							
	<a href="#">vllm</a>	Framework for LLM inference and serving							
	<a href="#">llmware</a>	LLM Development framework							
	<a href="#">Intel Extension for Transformers</a>	Allows for running models efficiently on Intel systems							
	<a href="#">HuggingFace PEFT</a>	Tool for finetuning models Supports a lot of parameter-efficient techniques (including QLoRA)							
	<a href="#">HuggingFace trl</a>	Training toolkit Supports RLHF (DPO, PPO)							
	<a href="#">Pinecone</a>	Vector database (recommended)							
	<a href="#">Qdrant</a>	Vector database							
	<a href="#">Weaviate</a>	Vector database							
Deployment	<a href="#">HuggingFace + Gradio</a>	Good choice for hosting custom models for inference as endpoints							
	<a href="#">SageMaker</a>	AWS platform for everything ML							
	<a href="#">streamlit</a>	Recommended for quick deployment and ease of use							
Miscellaneous	<a href="#">vast.ai</a>	Best on the market for renting GPUs							
	<a href="#">Paperspace</a>	Another good choice for renting GPUs							
	<a href="#">Lambda Labs</a>	Another good choice for renting GPUs							